

Introduction to Artificial Intelligence



VICTORIA UNIVERSITY OF
WELLINGTON
TE HERENGA WAKA

COMP307/AIML420

Neural Networks: Tutorial


Dr Andrew Lensen

Andrew.Lensen@vuw.ac.nz

NN Example: Your Turn!

- Calculate the outputs of this network (feedforward): to 2dp

I_1	I_2	w_{13}	w_{14}	w_{23}	w_{24}	w_{35}	w_{36}	w_{45}	w_{46}	b_3	b_4	b_5	b_6
0.90	-0.20	0.72	-0.31	0.10	-0.92	-0.37	0.43	-0.19	0.78	0.01	0.38	-0.13	0.78



z_3	
o_3	
z_4	
o_4	
z_5	
o_5	
z_6	
o_6	

Class = ____

Useful Formulae: Feedforward

- Weighted sum (ws) of a node:
$$z_j = \sum_i w_{ji} x_i + b_j$$
- Output of a node:
$$O_j = \varphi(z_j)$$
 - Where φ is the activation function.
- Assume φ is the sigmoid function:
$$\varphi(z_j) = \frac{1}{1 + e^{-z_j}}$$

NN Example: Your Turn!


- Calculate the new weights and biases (backprop): to 2dp

d_5	d_6	η	β_3	β_4	β_5	β_6
0	1	1				

Inputs

Hidden

Outputs




w_{13}	
w_{14}	
w_{23}	
w_{24}	
w_{35}	
w_{36}	
w_{45}	
w_{46}	
b_3	
b_4	
b_5	
b_6	

Useful Formulae: Backprop


- Error term of an output node: $\beta_j = d_j - O_j$
- Error term of a hidden node: $\beta_j = \sum_k w_{j \rightarrow k} O_k (1 - O_k) \beta_k$
 - (For the sigmoid activation function)
- Amount to change a weight: $\Delta w_{i \rightarrow j} = \eta O_i O_j (1 - O_j) \beta_j$
- Amount to change a bias: $\Delta b_j = \eta O_j (1 - O_j) \beta_j$

Notes on BP Algorithm

- *1 Epoch*: all input examples (entire training set, batch, ...)
- A target of 0 or 1 cannot reasonably be reached. Usually interpret an output > 0.9 or > 0.8 as '1'
- Training may require *thousands* of epochs. A convergence curve will help to decide when to stop (over-fitting?)



NNs for (Multi-Class) Classification



Training a Neural Network



- **Initialise** the weights (randomly)
- **Feedforward**
 - For each example, calculate the **predicted outputs** o_z using the current weights
 - Calculate the total **error** $\sum_z (d_z - o_z)^2$
- If the error is small enough, we can stop.
- Otherwise, we use **back propagation** to adjust the weights to make the error *smaller*.
 - Uses gradient descent (GD)

Back Propagation (BP) Algorithm

- Estimate the contribution (gradient) of each weight to the *error*, i.e. how much the error will be reduced by changing the weight (gradient)
- Change each weight (simultaneously) proportional to its contribution to reduce the error as much as possible
 - Move in the direction of the steepest gradient
- We calculate the contribution/gradient backwards (from the last/output layer to the first hidden layer)
- Error of a single output node is $d_z - o_z$
 - d_z means “*desired*”
 - o_z means “output” (i.e. what we actually got)

Back Propagation (BP) Algorithm

- How **big a change** should we make to **weight $w_{i \rightarrow j}$** ?
 - Make a **big change** if will improve error **a lot** (big contribution)
 - Make a **small change** if **little effect** on error (small contribution)



- β_j is how “**beneficial**” a change is for node j (“error term”)
- When changing $w_{i \rightarrow j}$, the error change should be:
 - Proportional to the **output**: o_i (larger output = more effect)
 - Proportional to the **slope of the activation function** at node j : slope_j
 - Proportional to error term of j (β_j)

BP Algorithm Implementation

- Let η be the learning rate (“eta”...)
- Initialise all weights (+bias) to small random values
- Until total error is small enough, repeat:
 - For each input example:
 - Feed forward pass to get predicted outputs
 - Compute $\beta_z = d_z - o_z$ for each output node
 - Compute $\beta_j = \sum_k w_{j \rightarrow k} o_k (1 - o_k) \beta_k$ for each hidden node (working backwards from last to first layer)
 - Compute (+store) the weight changes for all weights
$$\Delta w_{i \rightarrow j} = \eta o_i o_j (1 - o_j) \beta_j$$
(proportional to all 3 factors)
 - Sum up weight changes for all input examples
 - Change weights!